

Protein Structure Classification Using Machine Learning

Pin-Chieh Liao

María Victoria Liendro

Zahra Rahimi Pirkoohi

May 5, 2025

Abstract

Protein structure determination is essential for understanding protein function and facilitating applications in drug discovery and disease treatment. However, experimental methods for structure determination are time-consuming and expensive. We present a machine learning approach using bidirectional long short-term memory (BiLSTM) neural networks to predict the secondary structure of proteins from their amino acid sequences. Our model achieves 88% overall accuracy across all classes, providing a rapid and cost-effective alternative to experimental methods for initial structural characterization. This approach can accelerate research in structural biology by prioritizing proteins for experimental structure determination and providing preliminary insights into protein function.

1 Introduction

Proteins are essential biomolecules that serve as the workhorses of biological systems, performing a wide range of functions from enzymatic catalysis to structural support. The functionality of a protein is intrinsically linked to its three-dimensional structure, which is determined by the folding of its amino acid chain. Understanding protein structure is crucial for numerous applications, including drug discovery, disease mechanism elucidation, and biotechnology development.

However, experimental determination of protein structures through methods like X-ray crystallography, NMR spectroscopy, and cryo-electron microscopy is time-consuming and expensive. Due now, millions of proteins have been discovered, but experimental structure determination for a single protein can take months and cost hundreds of thousands of dollars, creating a significant bottleneck in biological research.

In this project, we address this challenge by developing a machine learning approach to predict the secondary structure of proteins at the residue level directly from their amino acid sequences. By mapping each amino acid to its corresponding secondary structure element—such as helix, strand, or coil—we can provide valuable insights into protein function and prioritize candidates for experimental structure determination, thereby accelerating research in structural biology and drug development.

Our approach utilizes a bidirectional long short-term memory (BiLSTM) neural network architecture, which is particularly well-suited for capturing sequential patterns in protein data. Our results demonstrate that this method achieves high classification accuracy across all structural classes, providing a fast and cost-effective alternative to experimental methods for initial structural characterization.

2 Problem Definition and Algorithm

2.1 Task Definition

We define our problem as follows: given a protein’s amino acid sequence, predict its structural class at residue-level according to the PDB (Protein Data Bank) database classification system. Each residue can be classified into:

- **B** – β -bridge (isolated β -bridge residue)
- **-** – Coil (random coil, unstructured)
- **E** – β -strand (extended sheet conformation)
- **G** – 3_{10} -helix (3-turn helix, tighter than α -helix)
- **H** – α -helix (standard right-handed helix)
- **I** – π -helix (wider than α -helix, rare)
- **S** – Bend (sharp bend in the backbone)
- **T** – Turn (reverse-turn, often hydrogen-bonded)

Input: A protein’s primary structure represented as a sequence of amino acids (a string of letters where each letter corresponds to one of the 20 standard amino acids).

Output: The predicted structural class of each amino acid.

This problem is crucial because quickly determining a protein’s structural class provides valuable insights into its potential function and can guide more targeted experimental investigations. For instance, membrane proteins frequently contain predominantly α -helical structures. In medical applications, understanding protein structure is crucial for drug design, as the binding sites for therapeutic compounds are determined by structural features.

In cancer treatment development, protein classification plays a vital role in targeted therapy, which focuses on specific proteins involved in cancer cell growth. The ability to rapidly classify newly discovered proteins can significantly accelerate the drug discovery pipeline, allowing researchers to develop compounds that precisely target cancer-related proteins while minimizing damage to healthy cells. This approach to medicine represents a significant advancement over traditional methods that often affect both healthy and cancerous cells.

2.2 Algorithm Definition

Our solution employs Bidirectional Long Short-Term Memory (BiLSTM) neural network, which is particularly effective for sequence data. The algorithm processes protein sequences as follows:

Preprocessing:

1. Tokenize the amino acid sequences and their corresponding secondary structures by mapping each unique element to an integer.
2. Pad sequences to a uniform length (600 amino acids) to have an uniform input.
3. Create one-hot encoded target vectors for the secondary structures.

Model Architecture:

1. Embedding layer: Transforms integer-encoded amino acids into dense vector representations (embedding dimension = 64)
2. BiLSTM layers: Two stacked bidirectional LSTM layers (hidden dimension = 128) that process the sequence both forward and backward
3. Dropout layers: Applied between LSTM layers to prevent overfitting (dropout rate = 0.3, recurrent dropout = 0.2)

Algorithm 1 Protein Structure Classification

```
1: function CLASSIFYPROTEINSTRUCTURE(amino_acid_sequence)
2:   tokenized_sequence  $\leftarrow$  tokenize(amino_acid_sequence)
3:   padded_sequence  $\leftarrow$  pad(tokenized_sequence, max_length=600)
4:   embedded_sequence  $\leftarrow$  embedding_layer(padded_sequence)
5:   hidden_states_1  $\leftarrow$  BiLSTM_layer_1(embedded_sequence)
6:   hidden_states_1  $\leftarrow$  dropout(hidden_states_1, rate=0.3)
7:   hidden_states_2  $\leftarrow$  BiLSTM_layer_2(hidden_states_1)
8:   hidden_states_2  $\leftarrow$  dropout(hidden_states_2, rate=0.3)
9:   class_probabilities  $\leftarrow$  softmax_layer(hidden_states_2)
10:  return argmax(class_probabilities)
11: end function
```

4. Dense output layer: A softmax layer that outputs probability distributions for each residue's secondary structure (output dimension = number of classes)

Example Walkthrough:

Consider the amino acid sequence, "KVPEAA".

1. This sequence is tokenized into integers using the amino acid dictionary (e.g., 'K'=1, 'V'=2, etc.)
2. The tokenized sequence is padded to reach our maximum length of 600
3. The embedding layer converts each token into a 64-dimensional vector
4. The BiLSTM layers process this sequence, capturing contextual patterns that relate to structural tendencies
5. The output layer produces probabilities for each amino acid: [0.01, 0.05, 0.05, 0.01, 0.77, 0.02, 0.03, 0.06] corresponding to [B, -, E, G, H, I, S, T]
6. Based on these probabilities, the first amino acid should be classified as 'H' an α -helix

Our approach treats protein sequences as a "language" where the 20 amino acids function as "words" that highly rely on context. The BiLSTM architecture is particularly well-suited for this task because it processes the sequence in both forward and backward directions, allowing it to learn longer-range dependencies that are crucial for protein structure prediction.

3 Experimental Evaluation

3.1 Experimental Setup

Datasets: We utilized protein sequence and structure data from the Protein Data Bank (PDB), which provides a repository of experimentally determined protein structures. For our specific experiment, we extracted 574 ubiquitin proteins with sequences shorter than 600 amino acids to address memory constraints. Ubiquitin was chosen as our focus because it is a relatively simple and well-studied protein family.

The distribution of our dataset shows a mean sequence length of 268.3 amino acids and a median of 242.0 amino acids. This distribution is representative of many functional proteins while avoiding extremely long sequences that would pose computational challenges.

In terms of amino acid distribution, our dataset shows leucine (L) as the most common amino acid, followed by lysine (K), glutamic acid (E), threonine (T), and isoleucine (I), see Figure 1.




Figure 1
Amino acid distribution




Figure 2
Secondary structure distribution

For secondary structure elements, our dataset contains helices (H) and extended strands (E) as the most common secondary structure types, see Figure 2.

The dataset was split into training (80%) and testing (20%) sets, ensuring proportional representation of each structural class in both sets.

Evaluation Metrics: We evaluated our model using the following metrics:

- Accuracy: The proportion of correctly classified amino acids
- Training and validation loss curves to assess model convergence
- Per-class prediction accuracy, which is correctly classified over samples count in each class, to identify strengths and weaknesses

Implementation Details:

- Framework: TensorFlow
- Training parameters: Batch size of 16, 25 epochs
- Optimization: Adam optimizer
- Loss function: Categorical cross-entropy
- Model variations: We tested multiple configurations with different embedding and hidden dimensions to find the optimal balance between performance and computational efficiency

3.2 Results

Our BiLSTM model achieved an overall accuracy of 88% on the test dataset, demonstrating strong performance across all structural classes. We evaluated several model configurations with varying embedding and hidden dimensions as shown in Table 1.

Embedding dimensions	Hidden dimensions	Training time (h)	Testing accuracy
16	32	~1.25	85.11%
32	64	~1.5	86.42%
64	128	>1.75	88.00%
128	256	—	Out of RAM

Table 1: Comparison of different model configurations

The training process showed consistent improvement over epochs, with validation accuracy plateauing around epoch 20, indicating sufficient training without significant overfitting. Figure 3 shows the training and validation accuracy curves over epochs.




Figure 3: Training and validation accuracy over epochs

Analysis of secondary structure prediction accuracy revealed that the model performs best on helices (H) with approximately 85% accuracy and extended strands (E) with approximately 78% accuracy, see Figure 4. This aligns with biochemical understanding, as these secondary structure elements have more distinctive amino acid composition patterns. It is worth noting that in our test set we did not have any π -helix, I, so it has zero percent of correct predictions in the plot.

When comparing BiLSTM against standard LSTM, we observed a significant performance advantage for BiLSTM, which achieved approximately 85% validation accuracy compared to only about 62% for LSTM at similar epochs, see Figure 5. This confirms our hypothesis that the bidirectional nature of BiLSTM provides crucial advantages for protein structure classification.




Figure 4: Percentage of correct prediction per secondary structure (blue) and the correctly classified count of each class (green)




Figure 5: Bidirectional LSTM and traditional LSTM accuracy

3.3 Discussion

The performance gap between BiLSTM and standard LSTM confirms the importance of processing protein sequences in both directions. This bidirectional approach allows the model to capture long-range interactions between amino acids that may be separated in the sequence but close in the folded three-dimensional structure.

We observed that increasing the model capacity through larger embedding and hidden dimensions improved performance, but with diminishing returns. The most significant limitation we encountered was hardware constraints, which prevented us from testing larger models with embedding dimensions above 64 and hidden dimensions above 128. This suggests that further improvements might be possible with more computational resources.

One limitation of our current approach is the constraint on sequence length. While our model handles sequences up to 600 amino acids effectively, many proteins in nature exceed this length. The longest known proteins can contain 33,000-37,000 amino acids. Handling such extremely long sequences would require alternative architectures or better computational resources.

4 Related Work

While many approaches to protein structure prediction exist, our work focuses specifically on the classification task rather than full 3D structure prediction. Traditional methods for protein structure prediction include:

- Homology modeling, which uses similar proteins with known structures as templates
- Physics-based approaches that simulate the folding process based on energy minimization principles
- Machine learning approaches ranging from support vector machines and random forests to deep learning architectures

Deep learning approaches have shown particular promise in recent years. Convolutional neural networks (CNNs) have been applied successfully to protein structure prediction tasks, particularly for local structure prediction. Recurrent neural networks (RNNs), especially LSTM variants, have demonstrated effectiveness in capturing sequential patterns in protein data.

Recent breakthroughs like AlphaFold by DeepMind have dramatically improved the accuracy of full 3D structure prediction, though these methods are computationally intensive and require substantial training data.

Our work differs from these approaches by focusing specifically on the classification task rather than full structure prediction, providing a faster alternative for applications where structural class information is sufficient. Additionally, our BiLSTM approach requires only the amino acid sequence as input, without the need for complex evolutionary information or multiple sequence alignments that many current state-of-the-art methods depend on.

5 Conclusion

Our work demonstrates that BiLSTM networks can effectively classify proteins into structural classes based solely on their amino acid sequences. With an overall accuracy of 88%, our approach provides a fast and cost-effective alternative to experimental methods for initial structural characterization.

Future work could focus on:

1. Incorporating additional sequence features, such as evolutionary information from multiple sequence alignments.
2. Creating an accessible web tool to make these predictions available to the broader scientific community
3. Exploring alternative architectures, such as transformer models, that may better handle very long protein sequences

By continuing to improve computational methods for protein structure analysis, we can help bridge the gap between the rapidly growing number of discovered protein sequences and our understanding of their structures and functions, ultimately accelerating advances in medicine and biotechnology.

References

- [1] Murzin, A. G., Brenner, S. E., Hubbard, T., & Chothia, C. (1995). SCOP: A structural classification of proteins database for the investigation of sequences and structures. *Journal of Molecular Biology*, 247(4), 536-540.
- [2] Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural Computation*, 9(8), 1735-1780.
- [3] Rost, B., & Sander, C. (1994). Combining evolutionary information and neural networks to predict protein secondary structure. *Proteins: Structure, Function, and Bioinformatics*, 19(1), 55-72.
- [4] Jumper, J., Evans, R., Pritzel, A., Green, T., Figurnov, M., Ronneberger, O., ... & Hassabis, D. (2021). Highly accurate protein structure prediction with AlphaFold. *Nature*, 596(7873), 583-589.
- [5] Yang, J., Yan, R., Roy, A., Xu, D., Poisson, J., & Zhang, Y. (2015). The I-TASSER Suite: protein structure and function prediction. *Nature Methods*, 12(1), 7-8.
- [6] Zhang, Y. (2008). Progress and challenges in protein structure prediction. *Current Opinion in Structural Biology*, 18(3), 342-348.
- [7] Cheng, J., Tegge, A. N., & Baldi, P. (2008). Machine learning methods for protein structure prediction. *IEEE Reviews in Biomedical Engineering*, 1, 41-49.

- [8] Wang, S., Peng, J., Ma, J., & Xu, J. (2016). Protein secondary structure prediction using deep convolutional neural fields. *Scientific Reports*, 6, 18962.
- [9] Heffernan, R., Yang, Y., Paliwal, K., & Zhou, Y. (2017). Capturing non-local interactions by long short-term memory bidirectional recurrent neural networks for improving prediction of protein secondary structure, backbone angles, contact numbers and solvent accessibility. *Bioinformatics*, 33(18), 2842-2849.
- [10] Wei, L., Ding, Y., Su, R., Tang, J., & Zou, Q. (2018). Prediction of human protein subcellular localization using deep learning. *Journal of Parallel and Distributed Computing*, 117, 212-217.